

# AUTOMATON AUDITOR

## Architecture Decision Rationale, Gap Analysis, and StateGraph Design

---

### I. Architecture Decision Rationale

The Automaton Auditor is intentionally architected as a **Digital Courtroom system** with layered reasoning:

- **Detective Layer** → Collects factual evidence only
- **Aggregation Layer** → Synchronizes parallel outputs
- **Judicial Layer** → Adversarial scoring via structured opinions
- **Chief Justice** → Deterministic synthesis and resolution

Each major technical decision was chosen to prevent specific architectural failure modes.

---

#### 1. Typed State: Why Pydantic + TypedDict + Reducers (Not Plain Dicts)

##### Current Implementation

- Evidence and JudicialOpinion are Pydantic models.
- AgentState is a TypedDict with explicit reducers:
  - operator.ior for evidences
  - operator.add for opinions

##### Why This Was Chosen

Using plain dictionaries would introduce several structural risks in a parallel LangGraph execution environment.

##### Failure Modes Prevented

1. **Silent schema drift**
  - Plain dicts allow accidental keys or malformed structures.
  - Pydantic enforces strict validation of:
    - score
    - criterion\_id

- judge
    - cited\_evidence
- 2. Parallel state corruption**
- LangGraph merges branches during fan-in.
  - Without reducers:
    - One branch could overwrite another.
  - With reducers:
    - opinions are appended safely.
    - evidences are merged safely by bucket.
- 3. Malformed LLM output contamination**
- Judges return JSON.
  - Pydantic validation ensures malformed outputs do not enter system state.

## Alternatives Considered

- Plain `dict` discipline → too fragile in multi-agent setting.
- Dataclasses → no built-in validation.
- Marshmallow → unnecessary complexity for internal validation.

**Conclusion:** Typed state is necessary for deterministic, parallel-safe orchestration.

---

## 2. AST Parsing vs Regex for Repository Analysis

### Current Implementation

- Python `ast` module is used to inspect:
  - `src/graph.py`
  - `src/state.py`

### Why AST Was Selected

Regex parsing is text-based and unaware of Python structure.

### Failure Modes Prevented

#### 1. Multiline definition breakage

Regex fails on:

```
2. class AgentState(
3.     TypedDict,
4.     total=False
5. ):
```

AST parses structure correctly.

6. **Nested definitions**  
Regex cannot reliably detect nested classes/functions.
7. **False positives**  
Regex may match commented code or docstrings.
8. **Structure misinterpretation**  
AST reads syntax trees, not raw strings.

## Alternatives Rejected

- Regex (too brittle)
- Executing inspected code (unsafe)
- Static analyzers (overkill for scope)

**Conclusion:** AST guarantees structural correctness over superficial pattern matching.

---

## 3. Sandboxing Strategy for Cloning Unknown Repositories

### Current Implementation

- `tempfile` directories
- No `os.system` usage
- Static inspection only
- No execution of cloned code

### Failure Modes Prevented

1. **Arbitrary code execution**  
No cloned code is executed.
2. **Shell injection**  
No direct shell calls.
3. **Filesystem contamination**  
Temporary directory isolation.
4. **Persistent malicious artifacts**  
Temporary directories are disposable.

### Alternatives Considered

- Docker isolation → excessive overhead.
- Running repo code → unsafe.

**Conclusion:** Lightweight sandboxing aligns with static audit design.

---

## 4. PDF Ingestion: RAG-lite Approach

### Current Implementation

- PDF chunking via PyPDF2
- In-memory lexical query
- No embeddings
- No external vector database

### Why This Was Chosen

The rubric requires presence/absence verification of conceptual terms.

### Failure Modes Prevented

1. **Dependency explosion**  
No vector DB required.
2. **Embedding variance**  
Lexical matching is deterministic.
3. **Cloud cost / API reliance**  
Entirely local.

### Trade-Off

- Lower semantic nuance than embeddings.
- Sufficient for rubric keyword validation.

**Conclusion:** RAG-lite matches scope without overengineering.

---

## 5. LLM Provider for Judges: Local DeepSeek via Ollama

### Current Implementation

- ChatOllama
- deepseek-r1:8b
- JSON hardening
- <think> stripping
- Structured fallback (neutral score)

### Why Local LLM

### Failure Modes Prevented

1. Cloud rate limits (429)
2. Quota exhaustion
3. Network outages
4. Billing unpredictability

## Trade-Off

- Occasional JSON formatting instability
- Slightly slower than high-end cloud APIs

Fallback strategy ensures the system never crashes.

**Conclusion:** Local LLM prioritizes reproducibility and stability.

---

## II. Gap Analysis and Forward Plan

This section explicitly identifies what is **not yet implemented** and outlines a concrete plan.

---

### Current Gaps

#### 1. Persona Divergence Enforcement

Currently:

- Personas differ by prompt instructions.
- No enforcement ensures divergence in reasoning style.

**Planned enhancement, not yet implemented:**

- Add divergence scoring heuristics.
  - Detect similarity in arguments across judges.
  - Force re-evaluation if opinions converge too closely.
- 

#### 2. Variance-Based Re-Evaluation in Synthesis

Currently:

- ChiefJustice performs deterministic weighted scoring.

- No variance threshold logic.

**Planned enhancement, not yet implemented:**

- Compute score variance across judges.
- If variance exceeds threshold:
  - Trigger additional deliberation pass.
  - Or increase TechLead weighting.

---

### 3. Confidence-Weighted Judging

Currently:

- Judges do not emit confidence values.

**Planned enhancement, not yet implemented:**

- Extend `JudicialOpinion` with confidence.
- Weight final score by confidence.

---

### 4. JSON Stability Hardening

Occasional fallback occurs due to malformed JSON.

**Planned enhancement, not yet implemented:**

- Enable Ollama JSON mode if supported.
- Add retry-on-parse-failure.
- Add stricter structural validation pre-Pydantic.

---

## Forward Plan (Sequenced)

### Phase 1 — Judicial Robustness

1. Enforce JSON mode.
2. Add retry logic.
3. Add citation validation.

## **Phase 2 — Dialectical Intelligence**

1. Implement variance-based escalation.
2. Introduce dissent analysis summary.
3. Detect persona collapse.

## **Phase 3 — Observability**


1. Log fallback frequency.
2. Track divergence metrics.
3. Track structured output violations.

Each phase is implementable independently.

---

## **III. StateGraph Architecture Diagram (Actual Implementation)**

This diagram reflects only currently implemented nodes.



## Data Flow

- Detectives → Evidence
  - Aggregator → Dict[str, List[Evidence]]
  - Judges → List[JudicialOpinion]
  - ChiefJustice → final\_report: str
-

## Error Handling (Currently Implemented)

- Repo clone failure → Evidence bucket with failure note.
  - Missing PDF → Factual absence.
  - Judge JSON failure → Neutral fallback opinions.
  - Invalid rubric → Synthetic `rubric_load_failed`.
-